

```
In [ ]: # Install TensorFlow
# !pip install -q tensorflow-gpu==2.0.0-beta1

try:
    %tensorflow_version 2.x # Colab only.
except Exception:
    pass

import tensorflow as tf
print(tf.__version__)
```

`%tensorflow\_version` only switches the major version: `1.x` or `2.x`.  
You set: `2.x` # Colab only.`. This will be interpreted as: `2.x`.

TensorFlow 2.x selected.  
2.1.0-rc1

```
In [ ]: # additional imports

import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras.layers import Input, Conv2D, Dense, Flatten, Dropout
from tensorflow.keras.models import Model
```

```
In [ ]: # Load in the data
fashion_mnist = tf.keras.datasets.fashion_mnist

(x_train, y_train), (x_test, y_test) = fashion_mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0
print("x_train.shape:", x_train.shape)
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-labels-idx1-ubyte.gz>  
32768/29515 [=====] - 0s 0us/step  
Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-images-idx3-ubyte.gz>  
26427392/26421880 [=====] - 0s 0us/step  
Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-labels-idx1-ubyte.gz>  
8192/5148 [=====] - 0s 0us/step  
Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-images-idx3-ubyte.gz>  
4423680/4422102 [=====] - 0s 0us/step  
x\_train.shape: (60000, 28, 28)

```
In [ ]: # the data is only 2D!
# convolution expects height x width x color
x_train = np.expand_dims(x_train, -1)
x_test = np.expand_dims(x_test, -1)
print(x_train.shape)
```

(60000, 28, 28, 1)

```
In [ ]: # number of classes
K = len(set(y_train))
```

```
print("number of classes:", K)
```

number of classes: 10

In [ ]:

```
# Build the model using the functional API
i = Input(shape=x_train[0].shape)
x = Conv2D(32, (3, 3), strides=2, activation='relu')(i)
x = Conv2D(64, (3, 3), strides=2, activation='relu')(x)
x = Conv2D(128, (3, 3), strides=2, activation='relu')(x)
x = Flatten()(x)
x = Dropout(0.2)(x)
x = Dense(512, activation='relu')(x)
x = Dropout(0.2)(x)
x = Dense(K, activation='softmax')(x)

model = Model(i, x)
```

In [ ]:

```
# Compile and fit
# Note: make sure you are using the GPU for this!
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
r = model.fit(x_train, y_train, validation_data=(x_test, y_test), epochs=15)
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/15

60000/60000 [=====] - 11s 175us/sample - loss: 0.5238 - accuracy: 0.8055 - val\_loss: 0.3963 - val\_accuracy: 0.8485

Epoch 2/15

60000/60000 [=====] - 6s 107us/sample - loss: 0.3679 - accuracy: 0.8616 - val\_loss: 0.3406 - val\_accuracy: 0.8744

Epoch 3/15

60000/60000 [=====] - 6s 107us/sample - loss: 0.3158 - accuracy: 0.8802 - val\_loss: 0.3336 - val\_accuracy: 0.8733

Epoch 4/15

60000/60000 [=====] - 6s 107us/sample - loss: 0.2813 - accuracy: 0.8942 - val\_loss: 0.3047 - val\_accuracy: 0.8878

Epoch 5/15

60000/60000 [=====] - 7s 109us/sample - loss: 0.2563 - accuracy: 0.9026 - val\_loss: 0.3044 - val\_accuracy: 0.8916

Epoch 6/15

60000/60000 [=====] - 7s 110us/sample - loss: 0.2363 - accuracy: 0.9104 - val\_loss: 0.2950 - val\_accuracy: 0.8955

Epoch 7/15

60000/60000 [=====] - 7s 108us/sample - loss: 0.2166 - accuracy: 0.9181 - val\_loss: 0.2928 - val\_accuracy: 0.8985

Epoch 8/15

60000/60000 [=====] - 7s 111us/sample - loss: 0.1986 - accuracy: 0.9243 - val\_loss: 0.2932 - val\_accuracy: 0.9001

Epoch 9/15

60000/60000 [=====] - 7s 112us/sample - loss: 0.1863 - accuracy: 0.9294 - val\_loss: 0.3185 - val\_accuracy: 0.8923

Epoch 10/15

60000/60000 [=====] - 7s 109us/sample - loss: 0.1712 - accuracy: 0.9355 - val\_loss: 0.3151 - val\_accuracy: 0.8970

Epoch 11/15

60000/60000 [=====] - 6s 107us/sample - loss: 0.1582 - accuracy: 0.9398 - val\_loss: 0.3154 - val\_accuracy: 0.8998

Epoch 12/15


60000/60000 [=====] - 6s 108us/sample - loss: 0.1509 - accuracy: 0.9414 - val\_loss: 0.3235 - val\_accuracy: 0.8979

Epoch 13/15  
 60000/60000 [=====] - 6s 107us/sample - loss: 0.1412 - accuracy: 0.9461 - val\_loss: 0.3347 - val\_accuracy: 0.9013  
 Epoch 14/15  
 60000/60000 [=====] - 6s 107us/sample - loss: 0.1326 - accuracy: 0.9489 - val\_loss: 0.3652 - val\_accuracy: 0.8961  
 Epoch 15/15  
 60000/60000 [=====] - 6s 107us/sample - loss: 0.1266 - accuracy: 0.9520 - val\_loss: 0.3516 - val\_accuracy: 0.9001

In [ ]:

```
# Plot loss per iteration
import matplotlib.pyplot as plt
plt.plot(r.history['loss'], label='loss')
plt.plot(r.history['val_loss'], label='val_loss')
plt.legend()
```


Out[ ]: &lt;matplotlib.legend.Legend at 0x7f23402c8860&gt;



In [ ]:

```
# Plot accuracy per iteration
plt.plot(r.history['accuracy'], label='acc')
plt.plot(r.history['val_accuracy'], label='val_acc')
plt.legend()
```

Out[ ]: &lt;matplotlib.legend.Legend at 0x7f23243183c8&gt;



In [ ]:

```

# Plot confusion matrix
from sklearn.metrics import confusion_matrix
import itertools

def plot_confusion_matrix(cm, classes,
                          normalize=False,
                          title='Confusion matrix',
                          cmap=plt.cm.Blues):
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """
    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
        print("Normalized confusion matrix")
    else:
        print('Confusion matrix, without normalization')

    print(cm)

    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=45)
    plt.yticks(tick_marks, classes)

    fmt = '.2f' if normalize else 'd'
    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, format(cm[i, j], fmt),
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")

    plt.tight_layout()
    plt.ylabel('True label')
    plt.xlabel('Predicted label')
    plt.show()

p_test = model.predict(x_test).argmax(axis=1)
cm = confusion_matrix(y_test, p_test)
plot_confusion_matrix(cm, list(range(10)))


```

Confusion matrix, without normalization

```

[[858  0 27 15  3  1 91  0  5  0]
 [ 1975  2 14  3  0  3  0  2  0]
 [ 13  1881 12 51  0 41  0  1  0]
 [ 12  6 20887 53  1 20  0  1  0]
 [ 0  1 74 10875  0 38  0  2  0]
 [ 0  0  1  0  0977  0 12  1  9]
 [115  0 87 26123  0641  0  8  0]
 [ 0  0  0  0  010  0962  0 28]
 [ 1  0  4  2  8  1  5  2975  2]
 [ 0  0  0  0  0  4  1 25  0970]]

```



```
In [ ]: # Label mapping
labels = '''T-shirt/top
Trouser
Pullover
Dress
Coat
Sandal
Shirt
Sneaker
Bag
Ankle boot'''.split("\n")
```

```
In [ ]: # Show some misclassified examples
misclassified_idx = np.where(p_test != y_test)[0]
i = np.random.choice(misclassified_idx)
plt.imshow(x_test[i].reshape(28,28), cmap='gray')
plt.title("True label: %s Predicted: %s" % (labels[y_test[i]], labels[p_test[i]]));
```

True label: Sneaker Predicted: Ankle boot

