

```
In [ ]:
try:
    %tensorflow_version 2.x # Colab only.
except Exception:
    pass

import tensorflow as tf
print(tf.__version__)
```

`%tensorflow_version` only switches the major version: `1.x` or `2.x`.
You set: `2.x` # Colab only.`. This will be interpreted as: `2.x`.

TensorFlow 2.x selected.
2.0.0-beta1

```
In [ ]:
# additional imports

import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras.layers import Input, Conv2D, Dense, Flatten, Dropout, GlobalMaxPo
from tensorflow.keras.models import Model
```

```
In [ ]:
# Load in the data
cifar10 = tf.keras.datasets.cifar10

(x_train, y_train), (x_test, y_test) = cifar10.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0
y_train, y_test = y_train.flatten(), y_test.flatten()
print("x_train.shape:", x_train.shape)
print("y_train.shape", y_train.shape)
```

x_train.shape: (50000, 32, 32, 3)
y_train.shape (50000,)

```
In [ ]:
# number of classes
K = len(set(y_train))
print("number of classes:", K)
```

number of classes: 10

```
In [ ]:
# Build the model using the functional API
def create_model():
    i = Input(shape=x_train[0].shape)

    x = Conv2D(32, (3, 3), activation='relu', padding='same')(i)
    x = BatchNormalization()(x)
    x = Conv2D(32, (3, 3), activation='relu', padding='same')(x)
    x = BatchNormalization()(x)
    x = MaxPooling2D((2, 2))(x)
    x = Conv2D(64, (3, 3), activation='relu', padding='same')(x)
    x = BatchNormalization()(x)
    x = Conv2D(64, (3, 3), activation='relu', padding='same')(x)
    x = BatchNormalization()(x)
    x = MaxPooling2D((2, 2))(x)
    x = Conv2D(128, (3, 3), activation='relu', padding='same')(x)
    x = BatchNormalization()(x)
```

```

x = Conv2D(128, (3, 3), activation='relu', padding='same')(x)
x = BatchNormalization()(x)
x = MaxPooling2D((2, 2))(x)

x = Flatten()(x)
x = Dropout(0.2)(x)
x = Dense(1024, activation='relu')(x)
x = Dropout(0.2)(x)
x = Dense(K, activation='softmax')(x)

model = Model(i, x)
return model

```

```

In [ ]: strategy = tf.distribute.MirroredStrategy()
# strategy = tf.distribute.experimental.CentralStorageStrategy()

```

```

In [ ]: print(f'Number of devices: {strategy.num_replicas_in_sync}')

```

Number of devices: 1

```

In [ ]: with strategy.scope():
        model = create_model()

        model.compile(loss='sparse_categorical_crossentropy',
                      optimizer='adam',
                      metrics=['accuracy'])

```

```

In [ ]: # Fit
r = model.fit(x_train, y_train, validation_data=(x_test, y_test), epochs=5)

```

WARNING: Logging before flag parsing goes to stderr.
W0811 18:03:54.175649 140429139875584 deprecation.py:323] From /tensorflow-2.0.0b1/python3.6/tensorflow/python/keras/layers/normalization.py:457: add_dispatch_support.<locals>.wrapper (from tensorflow.python.ops.array_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

Train on 1563 steps, validate on 313 steps

Epoch 1/15

1563/1563 [=====] - 23s 14ms/step - loss: 1.3045 - accuracy: 0.5541 - val_loss: 0.9532 - val_accuracy: 0.6640

Epoch 2/15

1563/1563 [=====] - 16s 10ms/step - loss: 0.8394 - accuracy: 0.7093 - val_loss: 0.8020 - val_accuracy: 0.7339

Epoch 3/15

1563/1563 [=====] - 16s 10ms/step - loss: 0.6854 - accuracy: 0.7643 - val_loss: 0.8026 - val_accuracy: 0.7286

Epoch 4/15

1563/1563 [=====] - 16s 10ms/step - loss: 0.5819 - accuracy: 0.8012 - val_loss: 0.6905 - val_accuracy: 0.7778

Epoch 5/15

1563/1563 [=====] - 16s 10ms/step - loss: 0.4859 - accuracy: 0.8322 - val_loss: 0.6967 - val_accuracy: 0.7767

Epoch 6/15

1563/1563 [=====] - 15s 10ms/step - loss: 0.4108 - accuracy: 0.8567 - val_loss: 0.6145 - val_accuracy: 0.8024

Epoch 7/15

```

1563/1563 [=====] - 15s 10ms/step - loss: 0.3517 - accuracy: 0.
8768 - val_loss: 0.6851 - val_accuracy: 0.7975
Epoch 8/15
1563/1563 [=====] - 16s 10ms/step - loss: 0.2934 - accuracy: 0.
8977 - val_loss: 0.6551 - val_accuracy: 0.8104
Epoch 9/15
1563/1563 [=====] - 16s 10ms/step - loss: 0.2536 - accuracy: 0.
9132 - val_loss: 0.6276 - val_accuracy: 0.8251
Epoch 10/15
1563/1563 [=====] - 16s 10ms/step - loss: 0.2205 - accuracy: 0.
9224 - val_loss: 0.6870 - val_accuracy: 0.8156
Epoch 11/15
1563/1563 [=====] - 16s 10ms/step - loss: 0.1921 - accuracy: 0.
9327 - val_loss: 0.7153 - val_accuracy: 0.8161
Epoch 12/15
1563/1563 [=====] - 15s 10ms/step - loss: 0.1732 - accuracy: 0.
9413 - val_loss: 0.6962 - val_accuracy: 0.8218
Epoch 13/15
1563/1563 [=====] - 16s 10ms/step - loss: 0.1587 - accuracy: 0.
9458 - val_loss: 0.7655 - val_accuracy: 0.8200
Epoch 14/15
1563/1563 [=====] - 16s 10ms/step - loss: 0.1512 - accuracy: 0.
9505 - val_loss: 0.8894 - val_accuracy: 0.8010
Epoch 15/15
1556/1563 [=====>.] - ETA: 0s - loss: 0.1332 - accuracy: 0.9550
W0811 18:07:53.185556 140432385251200 training_arrays.py:309] Your dataset ran out of da
ta; interrupting training. Make sure that your dataset can generate at least `steps_per_
epoch * epochs` batches (in this case, 23445 batches). You may need to use the repeat()
function when building your dataset.

```

```
In [ ]: 50000/391
```

```
Out[ ]: 127.8772378516624
```

```
In [ ]: 10000/79
```

```
Out[ ]: 126.58227848101266
```

```
In [ ]: # Compare this to non-distributed training
model2 = create_model()
model2.compile(loss='sparse_categorical_crossentropy',
               optimizer='adam',
               metrics=['accuracy'])
r = model2.fit(x_train, y_train, validation_data=(x_test, y_test), epochs=5)
```

Train on 50000 samples, validate on 10000 samples

```

Epoch 1/5
50000/50000 [=====] - 25s 502us/sample - loss: 1.3957 - accurac
y: 0.5369 - val_loss: 1.0386 - val_accuracy: 0.6328
Epoch 2/5
50000/50000 [=====] - 23s 465us/sample - loss: 0.8563 - accurac
y: 0.7029 - val_loss: 0.8461 - val_accuracy: 0.7063
Epoch 3/5
50000/50000 [=====] - 24s 471us/sample - loss: 0.7013 - accurac
y: 0.7592 - val_loss: 0.8004 - val_accuracy: 0.7334
Epoch 4/5
50000/50000 [=====] - 24s 473us/sample - loss: 0.5897 - accurac
y: 0.7959 - val_loss: 0.6652 - val_accuracy: 0.7788
Epoch 5/5

```

```
50000/50000 [=====] - 24s 471us/sample - loss: 0.5054 - accurac  
y: 0.8267 - val_loss: 0.6899 - val_accuracy: 0.7758
```