

Tensorflow 2.0 Variables

This notebook will teach you about variables and how to build a basic computation graph.

We will perform gradient descent manually to optimize a simple function.

If you're familiar with Tensorflow 1.x, you will find this useful as an example to demonstrate how we can do the same operations, but without sessions, initializers, etc.

```
In [ ]: # Install TensorFlow
# !pip install -q tensorflow-gpu==2.0.0-beta1

try:
    %tensorflow_version 2.x # Colab only.
except Exception:
    pass

import tensorflow as tf
print(tf.__version__)
```

`%tensorflow_version` only switches the major version: `1.x` or `2.x`.
You set: `2.x` # Colab only.`. This will be interpreted as: `2.x`.

TensorFlow 2.x selected.
2.0.0-beta1

```
In [ ]: # First, what is the difference between mutable and immutable?

# A tuple is immutable
# This should result in an error
a = (1,2,3)
a[0] = 5
```

```
-----
TypeError                                Traceback (most recent call last)
<ipython-input-2-0e96911569be> in <module>()
      1 a = (1,2,3)
----> 2 a[0] = 5
```

TypeError: 'tuple' object does not support item assignment

```
In [ ]: # A list is mutable
a = [1,2,3]
a[0] = 5
print(a)
```

[5, 2, 3]

```
In [ ]: # Now Tensorflow variables
a = tf.Variable(5.)
b = tf.Variable(3.)
print(a * b)
```

```
# Eager execution! No need for session.run() or variable initializer
```

```
tf.Tensor(15.0, shape=(), dtype=float32)
```

```
In [ ]: # Because it's a variable, it can be updated
a = a + 1
print(a)
```

```
tf.Tensor(6.0, shape=(), dtype=float32)
```

```
In [ ]: # Variables and constants
c = tf.constant(4.)
print(a * b + c)
```

```
tf.Tensor(22.0, shape=(), dtype=float32)
```

```
In [ ]: # Let's demonstrate a simple optimization problem
# L(w) = w**2

w = tf.Variable(5.)

# Now, let us define a loss function
def get_loss(w):
    return w ** 2

# Use "gradient tape" to record the gradients
def get_grad(w):
    with tf.GradientTape() as tape:
        L = get_loss(w)

    # Get the gradient
    g = tape.gradient(L, w)
    return g

# Define an optimizer
optimizer = tf.keras.optimizers.SGD(learning_rate=0.1)

# Store the losses
losses = []

# Perform gradient descent
for i in range(50):
    g = get_grad(w)
    optimizer.apply_gradients(zip([g], [w]))
    losses.append(get_loss(w))
```

WARNING: Logging before flag parsing goes to stderr.


W0812 20:20:03.417810 140436135712640 deprecation.py:323] From /tensorflow-2.0.0b1/python/n3.6/tensorflow/python/ops/math_grad.py:1205: add_dispatch_support.<locals>.wrapper (from tensorflow.python.ops.array_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

```
In [ ]: import matplotlib.pyplot as plt
plt.plot(losses)
print(f"Final loss: {get_loss(w)}")
```

Final loss: 5.0925916816879635e-09



```
In [ ]: # Let's do the same thing again, but manually

w = tf.Variable(5.)

# Store the losses
losses2 = []

# Perform gradient descent
for i in range(50):
    # This is doing: w = w - 0.1 * 2 * w
    # But we don't want to create a new Tensor
    w.assign(w - 0.1 * 2 * w)
    losses2.append(w ** 2)
```

```
In [ ]: plt.plot(losses, label="losses tf")
plt.plot(losses2, label="losses manual")
plt.legend()
```

Out[]: <matplotlib.legend.Legend at 0x7fb994fd99e8>

