

```
In [ ]: # Install TensorFlow
# !pip install -q tensorflow-gpu==2.0.0-beta1

try:
    %tensorflow_version 2.x # Colab only.
except Exception:
    pass

import tensorflow as tf
print(tf.__version__)
```

2.0.0-beta1

```
In [ ]: # additional imports

import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras.layers import Input, Conv2D, Dense, Flatten, Dropout, GlobalMaxPo
from tensorflow.keras.models import Model
```

```
In [ ]: # Load in the data
cifar10 = tf.keras.datasets.cifar10

(x_train, y_train), (x_test, y_test) = cifar10.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0
y_train, y_test = y_train.flatten(), y_test.flatten()
print("x_train.shape:", x_train.shape)
print("y_train.shape", y_train.shape)
```

x\_train.shape: (50000, 32, 32, 3)  
y\_train.shape (50000,)

```
In [ ]: # number of classes
K = len(set(y_train))
print("number of classes:", K)
```

number of classes: 10

```
In [ ]: # Build the model using the functional API
i = Input(shape=x_train[0].shape)
x = Conv2D(32, (3, 3), strides=2, activation='relu')(i)
x = Conv2D(64, (3, 3), strides=2, activation='relu')(x)
x = Conv2D(128, (3, 3), strides=2, activation='relu')(x)
x = Flatten()(x)
x = Dropout(0.5)(x)
x = Dense(1024, activation='relu')(x)
x = Dropout(0.2)(x)
x = Dense(K, activation='softmax')(x)

model = Model(i, x)
```

```
In [ ]: # Compile and fit
# Note: make sure you are using the GPU for this!
model.compile(optimizer='adam',
```

```

        loss='sparse_categorical_crossentropy',
        metrics=['accuracy'])
r = model.fit(x_train, y_train, validation_data=(x_test, y_test), epochs=15)

```

Train on 50000 samples, validate on 10000 samples

```

Epoch 1/15
50000/50000 [=====] - 17s 336us/sample - loss: 1.5736 - accurac
y: 0.4255 - val_loss: 1.3048 - val_accuracy: 0.5317
Epoch 2/15
50000/50000 [=====] - 16s 320us/sample - loss: 1.2732 - accurac
y: 0.5435 - val_loss: 1.1139 - val_accuracy: 0.6004
Epoch 3/15
50000/50000 [=====] - 16s 321us/sample - loss: 1.1454 - accurac
y: 0.5909 - val_loss: 1.0765 - val_accuracy: 0.6176
Epoch 4/15
50000/50000 [=====] - 16s 325us/sample - loss: 1.0497 - accurac
y: 0.6259 - val_loss: 0.9637 - val_accuracy: 0.6603
Epoch 5/15
50000/50000 [=====] - 16s 325us/sample - loss: 0.9764 - accurac
y: 0.6546 - val_loss: 0.9363 - val_accuracy: 0.6728
Epoch 6/15
50000/50000 [=====] - 16s 323us/sample - loss: 0.9164 - accurac
y: 0.6735 - val_loss: 0.9153 - val_accuracy: 0.6772
Epoch 7/15
50000/50000 [=====] - 16s 322us/sample - loss: 0.8660 - accurac
y: 0.6944 - val_loss: 0.8996 - val_accuracy: 0.6818
Epoch 8/15
50000/50000 [=====] - 17s 330us/sample - loss: 0.8213 - accurac
y: 0.7083 - val_loss: 0.8806 - val_accuracy: 0.6954
Epoch 9/15
50000/50000 [=====] - 16s 328us/sample - loss: 0.7807 - accurac
y: 0.7233 - val_loss: 0.8833 - val_accuracy: 0.6951
Epoch 10/15
50000/50000 [=====] - 16s 327us/sample - loss: 0.7470 - accurac
y: 0.7359 - val_loss: 0.8368 - val_accuracy: 0.7058
Epoch 11/15
50000/50000 [=====] - 16s 325us/sample - loss: 0.7121 - accurac
y: 0.7462 - val_loss: 0.8376 - val_accuracy: 0.7092
Epoch 12/15
50000/50000 [=====] - 17s 330us/sample - loss: 0.6808 - accurac
y: 0.7560 - val_loss: 0.8430 - val_accuracy: 0.7115
Epoch 13/15
50000/50000 [=====] - 16s 328us/sample - loss: 0.6576 - accurac
y: 0.7671 - val_loss: 0.8284 - val_accuracy: 0.7110
Epoch 14/15
50000/50000 [=====] - 16s 328us/sample - loss: 0.6402 - accurac
y: 0.7702 - val_loss: 0.8487 - val_accuracy: 0.7067
Epoch 15/15
50000/50000 [=====] - 16s 328us/sample - loss: 0.6232 - accurac
y: 0.7769 - val_loss: 0.8320 - val_accuracy: 0.7136

```


In [ ]:

```

# Plot loss per iteration
import matplotlib.pyplot as plt
plt.plot(r.history['loss'], label='loss')
plt.plot(r.history['val_loss'], label='val_loss')
plt.legend()


```

Out[ ]: <matplotlib.legend.Legend at 0x7f04d6f0db00>



```
In [ ]: # Plot accuracy per iteration
plt.plot(r.history['accuracy'], label='acc')
plt.plot(r.history['val_accuracy'], label='val_acc')
plt.legend()
```

Out[ ]: <matplotlib.legend.Legend at 0x7f04d6b34208>



```
In [ ]: # Plot confusion matrix
from sklearn.metrics import confusion_matrix
import itertools

def plot_confusion_matrix(cm, classes,
                           normalize=False,
                           title='Confusion matrix',
                           cmap=plt.cm.Blues):
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """
    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
        print("Normalized confusion matrix")
    else:
        print('Confusion matrix, without normalization')
```

```

print(cm)

plt.imshow(cm, interpolation='nearest', cmap=cmap)
plt.title(title)
plt.colorbar()
tick_marks = np.arange(len(classes))
plt.xticks(tick_marks, classes, rotation=45)
plt.yticks(tick_marks, classes)

fmt = '.2f' if normalize else 'd'
thresh = cm.max() / 2.
for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
    plt.text(j, i, format(cm[i, j], fmt),
             horizontalalignment="center",
             color="white" if cm[i, j] > thresh else "black")

plt.tight_layout()
plt.ylabel('True label')
plt.xlabel('Predicted label')
plt.show()

p_test = model.predict(x_test).argmax(axis=1)
cm = confusion_matrix(y_test, p_test)
plot_confusion_matrix(cm, list(range(10)))


```

Confusion matrix, without normalization

```

[[741 16 36 26 12 2 6 11 113 37]
 [ 14 781 9 9 2 4 6 2 47 126]
 [ 62 9 545 91 102 66 50 29 28 18]
 [ 20 9 51 548 51 171 58 33 30 29]
 [ 40 10 62 79 617 43 47 78 19 5]
 [ 15 2 52 214 38 572 13 55 21 18]
 [ 6 9 45 101 43 44 711 7 18 16]
 [ 24 3 38 48 59 64 5 726 6 27]
 [ 49 30 9 17 6 3 1 3 865 17]
 [ 35 105 16 19 4 7 1 11 46 756]]

```



In [ ]:

```

# Label mapping
labels = ''airplane
         automobile
         bird

```

```
cat
deer
dog
frog
horse
ship
truck''.split()
```

In [ ]:

```
# Show some misclassified examples
# TODO: add label names
misclassified_idx = np.where(p_test != y_test)[0]
i = np.random.choice(misclassified_idx)
plt.imshow(x_test[i], cmap='gray')
plt.title("True label: %s Predicted: %s" % (labels[y_test[i]], labels[p_test[i]]));
```

